

nesC 1.2 Language Reference Manual

David Gay, Philip Levis, David Culler, Eric Brewer

July 2004

1 Introduction

nesC is an extension to C [2] designed to embody the structuring concepts and execution model of TinyOS [1]. TinyOS is an event-driven operating system designed for sensor network nodes that have very limited resources (e.g., 8K bytes of program memory, 512 bytes of RAM). TinyOS has been reimplemented in nesC. This manual describes v1.2 of nesC, changes from v1.0 and v1.1 are summarised in Section 2.

The basic concepts behind nesC are:

- Separation of construction and composition: programs are built out of *components*, which are assembled (“wired”) to form whole programs. Components define two scopes, one for their specification (containing the names of their *interfaces*) and one for their implementation. Components have internal concurrency in the form of *tasks*. Threads of control may pass into a component through its interfaces. These threads are rooted either in a task or a hardware interrupt.
- Specification of component behaviour in terms of set of *interfaces*. Interfaces may be provided or used by the component. The provided interfaces are intended to represent the functionality that the component provides to its user, the used interfaces represent the functionality the component needs to perform its job.
- Interfaces are bidirectional: they specify a set of functions to be implemented by the interface’s provider (*commands*) and a set to be implemented by the interface’s user (*events*). This allows a single interface to represent a complex interaction between components (e.g., registration of interest in some event, followed by a callback when that event happens). This is critical because all lengthy commands in TinyOS (e.g. send packet) are non-blocking; their completion is signaled through an event (send packet done). The interface forces a component that calls the “send packet” command to provide an implementation for the “send packet done” event.

Typically commands call “downwards”, i.e., from application components to those closer to the hardware, while events call “upwards”. Certain primitive events are bound to hardware interrupts (the nature of this binding is system-dependent, so is not described further in this reference manual).

- Components are statically linked to each other via their interfaces. This increases runtime efficiency, encourages robust design, and allows for better static analysis of programs.
- nesC is designed under the expectation that code will be generated by whole-program compilers. This allows for better code generation and analysis. An example of this is nesC's compile-time data race detector.
- The concurrency model of nesC is based on run-to-completion tasks, and interrupt handlers which may interrupt tasks and each other. The nesC compiler signals the potential data races caused by the interrupt handlers.

This document is a reference manual for nesC rather than a tutorial. The TinyOS tutorial¹ presents a gentler introduction to nesC.

The rest of this document is structured as follows: Section 2 summarises the new features in nesC since v1.0. Section 3 presents the notation used in the reference manual, and Section 4 the scoping and naming rules of nesC. Sections 5 and 6 present interfaces and components, while Sections 7 and 8 explain how components are implemented. Section 9 presents nesC's concurrency model and data-race detection. Section 10 explains how C files, nesC interfaces and components are assembled into an application. Section 11 covers the remaining miscellaneous features of nesC. Finally, Appendix A fully defines nesC's grammar (as an extension to the C grammar from Appendix A of Kernighan and Ritchie (K&R) [2, pp234–239]), and Appendix B is a glossary of the terms used in this reference manual.

2 Changes

The changes from nesC 1.1 to 1.2 are:

- Generic interfaces: interfaces can now take type parameters (allowing, e.g., a single interface definition for a queue of any type of values).
- Generic components: components can now be instantiated (at compile-time), and can take value and type arguments (e.g., a generic queue component would take type and queue size arguments).

The changes from nesC 1.0 to 1.1 are:

1. **atomic** statements. These simplify implementation of concurrent data structures, and are understood by the new compile-time data-race detector.
2. Compile-time data-race detection gives warnings for variables that are potentially accessed concurrently by two interrupt handlers, or an interrupt handler and a task.
3. Commands and events which can safely be executed by interrupt handlers must be explicitly marked with the **async** storage class specifier.

¹Available with the TinyOS distribution at <http://webs.cs.berkeley.edu>

4. The results of calls to commands or events with “fan-out” are automatically combined by new type-specific combiner functions.
5. `uniqueCount` is a new *constant function* (Section 11.2) which counts uses of `unique`.
6. The NESC preprocessor symbol indicates the language version. It is 110 for nesC 1.1, 120 for nesC 1.2.

3 Notation

The `typewriter` font is used for nesC code and for filenames. Single symbols in italics, with optional subscripts, are used to refer to nesC entities, e.g., “component *K*” or “value *v*”.

The grammar of nesC is an extension the ANSI C grammar. We chose to base our presentation on the ANSI C grammar from Appendix A of Kernighan and Ritchie (K&R) [2, pp234–239]. We will not repeat productions from that grammar here. Words in *italics* are non-terminals and non-literal terminals, `typewriter` words and symbols are literal terminals. The subscript *opt* indicates optional terminals or non-terminals. In some cases, we change some ANSI C grammar rules. We indicate this as follows: *also* indicates additional productions for existing non-terminals, *replaced by* indicates replacement of an existing non-terminal.

Explanations of nesC constructs are presented along with the corresponding grammar fragments. In these fragments, we sometimes use `...` to represent elided productions (irrelevant to the construct at hand). Appendix A presents the full nesC grammar.

Several examples use the `uint8_t` and `uint16_t` types (from the C99 standard `inttypes.h` file) and the standard TinyOS `result_t` type (which represents success vs failure of an operation).

4 Scopes and Name Spaces in nesC

nesC includes the standard C name spaces: *object*, which includes variables, functions, typedefs, and enum-constants; *label* for `goto` labels; *tag* for `struct`, `union`, `enum` tags. It adds an additional *component* name space for component and interface definitions. For simplicity, we assume that each scope contains all four name spaces, though language restrictions mean that many of these name spaces are empty (e.g., all component and interface definitions are global, so the *component* name space is empty in all but the global scope).

nesC follows the standard C scoping rules, with the following additions:

- Each interface definition introduces two scopes. The *interface parameter scope* is nested in the global scope and contains the parameters of generic interface definitions. The *interface scope* is nested in the interface parameter scope and contains the interface’s commands and events.
- Each component definition introduces three new scopes. The *component parameter scope* is nested in the global scope and contains the parameters of generic component definitions. The *specification scope* is nested in the component parameter scope and contains the component’s specification elements.

The *implementation scope* is nested in the specification scope. For configurations, the implementation scope contains the names by which this component refers to its included components (Section 8.1). For modules, the implementation scope holds the tasks, C declarations and definitions that form the module's body. These declarations, etc may introduce their own nested scopes within the implementation scope, following the usual C scoping rules.

As usual in C, scopes must not have multiple definitions of the same name within the same name space.

5 Interface and Component Specification

A nesC *interface definition* specifies a bi-directional interaction between two components, known as the *provider* and *user*. Interactions via interfaces are specified by two sets of functions: *commands* are function calls from the user to the provider component, *events* are function calls from the provider to the user component. In many cases, the provider component is providing some service (e.g., sending messages over the radio) and commands represent requests, events responses.

An interface definition has a unique name, optional C type parameters, and contains declarations for its command and event functions. An interface definition with type parameters is called a *generic interface definition*.

An *interface type* is a reference to an interface definition and, if the referenced definition is generic, corresponding type arguments. Components can only be connected via two interfaces with the same type.

A component's *specification* is the set of interfaces that it provides and uses. Component specifications can also contain *bare* commands and events (i.e., not contained in an interface), to simplify the exposition we defer discussion of these to Section 5.4. Each provided or used interface has a name and an interface type.

For instance, the following source code

```
interface SendMsg { // send a radio message
    command result_t send(uint16_t address, uint8_t length, TOS_MsgPtr msg);
    event result_t sendDone(TOS_MsgPtr msg, result_t success);
}

interface Init<t> { // a generic interface definition
    command void doit(t x);
}

module Simple {
    provides interface Init<int> as MyInit;
    uses interface SendMsg as MyMessage;
} ...
```

shows two interface definitions, `SendMsg` and `Init`, and the specification of the `Simple` component. The specification of `Simple` has two elements: `MyInit`, a provided interface of type `Init<int>` and

`MyMessage` a used interface of type `SendMsg`. `Simple` must implement the `MyInit.doit` command and the `MyMessage.sendDone` event. It can call the `MyMessage.send` command.

The rest of this section covers interface definitions, interface types and component specifications in detail. The sections on component definition (Section 6) and implementations (Sections 7 and 8) explain how commands and events are called and implemented, and how components are linked together through their interfaces.

5.1 Interface Definitions

Interface definitions have the following syntax:

interface-definition:
`interface identifier type-parametersopt { declaration-list }`

Interface definitions have a name (*identifier*) with global scope. This name belongs to the component name space (Section 4), so interface definitions must have a name distinct from other interface definitions and from components, however they do not conflict with regular C declarations.

The *type-parameters* is a list of optional C type parameters for this interface definition:

type-parameters:
`< type-parameter-list >`

type-parameter-list:
`identifier`
`type-parameter-list , identifier`

These parameters belong to the object name space of the interface's parameter scope (Section 4) and are therefore visible in the *declaration-list*. See Section 11.3 for how type parameters interact with C's type system (in brief, these type parameters can be used like `typedef`'d types). An interface definition with type parameters is called a *generic interface definition*.

The *declaration-list* of an interface definition specifies a set of commands and events. It must consist of function declarations with the `command` or `event` storage class:

storage-class-specifier: also one of
`command event async`

The optional `async` keyword indicates that the command or event can be executed in an interrupt handler (see Section 9). The interface's commands and events belong to the object name space of the interface's scope (Section 4).

The example code above showed two simple interface definitions (`SendMsg` and `Init`). The following

```
interface Queue<t> {
    async command void push(t x);
    async command t pop();
    async command bool empty();
}
```

```

    async command bool full();
}

```

defines a generic interface `Queue` with a single type parameter, defining four commands which can be executed in an interrupt handler.

5.2 Interface Types

An interface type is specified by giving the name of an interface definition and, for generic interface definitions, any required type arguments:

interface-type:
interface *identifier* *type-arguments*_{opt}

type-arguments:
< *type-argument-list* >

type-argument-list:
type-name
type-argument-list , *type-name*

There must be as many types in *type-arguments* as there are parameters in the interface definition's type parameter list. Type arguments can not be incomplete or of function or array type.

Two interface types are the same if they refer to the same interface definition and their corresponding type arguments (if any) are of the same C type. Example interface types are `interface SendMsg` and `interface Queue<int>`.

5.3 Component Specification

The first part of a component's definition (see Section 6) is its *specification*, a declaration of provided or used specification elements, where each element is an interface (or, Section 5.4, a bare command or event):

component-specification:
{ *uses-provides-list* }

uses-provides-list:
uses-provides
uses-provides-list *uses-provides*

uses-provides:
uses *specification-element-list*
provides *specification-element-list*

specification-element-list:
specification-element

{ *specification-elements* }

specification-elements:

specification-element

specification-elements specification-element

There can be multiple **uses** and **provides** directives in a component specification. Multiple used or provided specification elements can be grouped in a single directive by surrounding them with { and }. For instance, these two specifications are identical:

```
module A1 {
    uses interface X;
    uses interface Y;
} ...

module A1 {
    uses {
        interface X;
        interface Y;
    }
} ...
```

An interface declaration has an interface type and an optional name:

specification-element:

interface-type instance-name_{opt} instance-parameters_{opt}

...

instance-name:

as *identifier*

instance-parameters:

[*parameter-type-list*]

If the name is omitted, the interface's name is the same as the name of the interface definition specified by the interface type: **interface** SendMsg means the same thing as **interface** SendMsg **as** SendMsg and **interface** Queue<int> is the same as **interface** Queue<int> **as** Queue. A specification can contain independent interfaces of the same interface type, e.g.,

```
provides interface X as X1;
uses interface X as X2;
```

The interface names belong to the object name space of the specification's scope (Section 4), thus there is no confusion between interface names and interface definition names (the latter are in the component name space).

An interface declaration without *instance-parameters* (e.g., **interface** X **as** Y) declares a *simple interface* to this component. A declaration with *instance-parameters* (e.g., **interface** SendMsg S[uint8_t id]) declares a *parameterised interface*, corresponding to multiple interfaces to this component, one for each distinct tuple of parameter values (so **interface** SendMsg **as** S[uint8_t

`id, uint8_t id2]` declares 256 * 256 interfaces of type `SendMsg`). The types of the *parameters* must be integral types (*enums* are not allowed at this time).

The specification for `AMStandard`, a component that dispatches messages received from the serial port and the radio to the application based on the “active message id” stored in the message, and sends messages to the radio or serial port depending on the selected destination address, is typical of many TinyOS system components:

```
module AMStandard {
  provides {
    interface StdControl;

    // The interface are parameterised by the active message id
    interface SendMsg[uint8_t id];
    interface ReceiveMsg[uint8_t id];
  }
  uses {
    interface StdControl as RadioControl;
    interface SendMsg as RadioSend;
    interface ReceiveMsg as RadioReceive;

    interface StdControl as SerialControl;
    interface SendMsg as SerialSend;
    interface ReceiveMsg as SerialReceive;
  }
} ...
```

It provides or uses nine interfaces:

- The provided interface `StdControl` of type `StdControl` supports initialisation of `AMStandard`.
- The provided parameterised interfaces of type `SendMsg` and `ReceiveMsg` (named `SendMsg` and `ReceiveMsg` respectively) support dispatching of received messages and sending of messages with a particular active message id
- The used interfaces control, send and receive messages from the radio and serial port respectively (another TinyOS component, the `GenericComm` configuration wires `AMStandard` to the lower-level components providing radio and serial port networking).

5.4 Bare Commands and Events

Commands or events can be included directly as specification elements by including a standard C function declaration with `command` or `event` as its storage class specifier:

```
specification-element:
    declaration
    ...
```


storage-class-specifier: also one of
`command event async`

It is a compile-time error if the *declaration* is not a function declaration with the `command` or `event` storage class. As in interfaces, `async` indicates that the command or event can be called from an interrupt handler. These bare command and events also belong to the object name space of the specification's scope (Section 4).

As with interface declarations, commands (events) can have instance parameters; these are placed before the function's regular parameter list, e.g., `command void send[uint8_t id](int x):`

direct-declarator: also
direct-declarator instance-parameters (parameter-type-list)
`...`

Commands are *simple commands* (*simple events*) if they have no instance parameters, or *parameterised commands* (*parameterised events*) if instance parameters are specified. Note that instance parameters are only allowed on commands or events within component specifications, not within interface definitions.

Module `M` of Figure 1 (Section 8.4) shows an example of a component specification with a bare command.

5.5 Command and Event Terminology

We say that a command (event) F provided in the specification of component K is *provided command (event) F of K* ; similarly, a command (event) used in the specification of component K is *used command (event) F of K* .

A command F in a provided interface X of component K is provided command $X.F$ of K ; a command F in a used interface X of K is used command $X.F$ of K ; an event F in a provided interface X of K is used event $X.F$ of K ; and an event F in a used interface X of K is provided event $X.F$ of K (note the reversal of used and provided for events due to the bidirectional nature of interfaces).

We use greek letters α, β, \dots to refer to any command or event of a component when the distinction between bare commands (events) and commands (events) in interfaces is not relevant. Commands or events α of K may be parameterised or simple, depending on the parameterised or simple status of the specification element to which they correspond.

We will often simply refer to the “command or event α of K ” when the used/provided distinction is not relevant.

6 Component Definition

A nesC component definition has a name, optional arguments, a specification and an implementation:

```

component:
    comp-kind identifier comp-parametersopt component-specification implementation

comp-kind: module
    configuration
    generic module
    generic configuration

implementation:
    module-implementation
    configuration-implementation

```

The component name belongs to the component name space of the global scope, hence must be distinct from other components and from interface definitions. There are two kinds of components: *modules* are components which are implemented with C code (Section 7), *configurations* are components which are implemented by assembling other components (Section 8). Components with parameters are called *generic components*, they must be instantiated in a configuration before they can be used (Section 8). Components without parameters exist as a single instance which is implicitly instantiated. The component's definition must reflect these distinctions (the *comp-kind* rule): for instance, a generic module A is defined with **generic module A()** {..., a non-generic configuration B is defined with **configuration B** {...

6.1 Generic Components

Generic component parameter lists are similar to function parameter lists, but allow for type parameters by (re)using the **typedef** keyword:

```

comp-parameters:
    ( component-parameter-list )

component-parameter-list:
    component-parameter
    component-parameter-list , component-parameter

component-parameter:
    parameter-declaration
    typedef identifier

```

The parameters belong to the object name space of the component's parameter scope (Section 4), and are hence visible both in the component's specification and implementation. Non-type parameters must be of arithmetic or **char[]** type. These parameters can be used as follows:

- Type parameters can be used as if the argument was of some unknown **typedef**'d type (for more details, see Section 11.3).

- Non-type parameters are constants of some unknown value (for more details, see Section 11.1); they can be used in any constant expression. They cannot be assigned to.

An instantiation with arguments a_1, \dots, a_n of generic component X with parameters p_1, \dots, p_n behaves like a new, non-generic component with the specification and implementation of X where all uses of parameter p_i have been replaced by the corresponding a_i value or type.² Section 10.2 details when generic components get instantiated.

6.2 Examples

Some examples (with simple specifications):

```
module A { provides interface X; } ...
generic configuration B() { uses interface Y; } ...
generic module AQueue(int n, typedef t) { provides interface Queue<t>; } ...
```

A is a simple module, B a generic configuration with no arguments but which can be instantiated multiple times, AQueue a generic module implementing an n entry queue with elements of type t . Note how AQueue instantiates the generic interface Queue with its type parameter t .

7 Modules

Modules implement a component specification with C code:

```
module-implementation:
    implementation { translation-unit }
```

where *translation-unit* is a list of C declarations and definitions (see K&R [2, pp234–239]).

The top-level declarations of the module's *translation-unit* belong to the module's implementation scope (Section 4). These declarations have indefinite extent and can be: any standard C declaration or definition, a task declaration or definition (placed in the object name space), a commands or event implementation.

7.1 Implementing the Module's Specification

The *translation-unit* must implement all provided commands (events) α of the module (i.e., all directly provided commands and events, all commands in provided interfaces and all events in used interfaces). A module can call any of its commands and signal any of its events.

These command and event implementations are specified with the following C syntax extensions:

²The most straightforward implementation of these semantics for generic modules is to duplicate X 's code. In some cases (e.g., no arguments to X), a nesC compiler might be able to share code between the instances of X at some runtime cost.

storage-class-specifier: also one of
`command event async`

declaration-specifiers: also
`default declaration-specifiers`

direct-declarator: also
`identifier . identifier`
`direct-declarator interface-parameters (parameter-type-list)`

The implementation of simple command or event α has the syntax of a C function definition for α (note the extension to *direct-declarator* to allow `.` in function names) with storage class `command` or `event`. Additionally, the `async` keyword must be included iff it was included in α 's declaration. For example, in a module that provides interface `Send` of type `SendMsg`:

```
command result_t Send.send(uint16_t address, uint8_t length, TOS_MsgPtr msg) {  
    ...  
    return SUCCESS;  
}
```

The implementation of parameterised command or event α with interface parameters P has the syntax of a C function definition for α with storage class `command` or `event` where the function's regular parameter list is prefixed with the parameters P within square brackets (this is the same syntax as parameterised command or event declarations within a component specification). These interface parameter declarations P belong to α 's function-parameter scope and have the same extent as regular function parameters. For example, in a module that provides interface `Send[uint8_t id]` of type `SendMsg`:

```
command result_t Send.send[uint8_t id](uint16_t address, uint8_t length,  
                                         TOS_MsgPtr msg) {  
    ...  
    return SUCCESS;  
}
```

Compile-time errors are reported when:

- There is no implementation for a provided command or event.
- The type signature, optional interface parameters and presence or absence of the `async` keyword of a command or event does not match that given in the module's specification.

7.2 Calling Commands and Signaling Events

The following extensions to C syntax are used to call events and signal commands:

postfix-expression:

```
postfix-expression [ argument-expression-list ]
call-kindopt primary ( argument-expression-listopt )
...
```

call-kind: one of

call signal post

A simple command α is called with `call α (...)`, a simple event α is signaled with `signal α (...)`. For instance, in a module that uses interface `Send` of type `SendMsg`: `call Send.send(1, sizeof(Message), &msg1)`.

A parameterised command α (respectively, an event) with n interface parameters of type τ_1, \dots, τ_n is called with interface parameters expressions e_1, \dots, e_n as follows: `call α [e_1, \dots, e_n](...)` (respectively, `signal α [e_1, \dots, e_n](...)`). Interface parameter expression e_i must be assignable to type τ_i ; the actual interface parameter value is e_i cast to type τ_i . For instance, in a module that uses interface `Send[uint8_t id]` of type `SendMsg`:

```
int x = ...;
call Send.send[x + 1](1, sizeof(Message), &msg1);
```

Execution of commands and events is immediate, i.e., `call` and `signal` behave similarly to function calls. The actual command or event implementations executed by a `call` or `signal` expression depend on the wiring statements in the program's configurations. These wiring statements may specify that 0, 1 or more implementations are to be executed. When more than 1 implementation is executed, we say that the module's command or event has "fan-out".

A module can specify a default implementation for a used command or event α that it calls or signals. A compile-time error occurs for default implementations of provided commands or events. Default implementations are executed when α is not connected to any command or event implementation. A default command or event is defined by prefixing a command or event implementation with the `default` keyword:

```
declaration-specifiers: also
    default declaration-specifiers
```

For instance, in a module that uses interface `Send` of type `SendMsg`:

```
default command result_t Send.send(uint16_t address, uint8_t length,
                                   TOS_MsgPtr msg) {
    return SUCCESS;
}
/* call is allowed even if interface Send is not connected */
... call Send.send(1, sizeof(Message), &msg1) ...
```

Section 8.4 specifies what command or event implementations are actually executed and what result gets returned by `call` and `signal` expressions.

7.3 Tasks

A task is an independent locus of control defined by a function of storage class `task` returning `void` and with no arguments: `task void myTask() { ... }`.³ A task can also have a forward declaration, e.g., `task void myTask();`.

Tasks are posted by prefixing a call to the task with `post`, e.g., `post myTask()`. `Post` returns immediately; its return value is 1 if the task was successfully posted for independent execution, 0 otherwise. The type of a post expression is `unsigned char`.

storage-class-specifier: also one of
`task`

call-kind: also one of
`post`

nesC’s concurrency model, including tasks, is presented in detail in Section 9.

7.4 Atomic statements

Atomic statements:

atomic-stmt:
`atomic statement`

guarantee that the statement is executed “as-if” no other computation occurred simultaneously. It is used to implement mutual exclusion, for updates to concurrent data structures, etc. A simple example is:

```
bool busy; // global

void f() {
    bool available;

    atomic {
        available = !busy;
        busy = TRUE;
    }
    if (available) do_something;
    atomic busy = FALSE;
}
```

Atomic sections should be short, though this is not currently enforced in any way. Control may only flow “normally” in or out of an atomic statement: any `goto`, `break` or `continue` that jumps

³nesC functions with no arguments are declared with `()`, not `(void)`. See Section 11.4.

in or out of an atomic statement is an error. A **return** statement is never allowed inside an atomic statement.

Section 9 discusses the relation between **atomic** and nesC’s concurrency model and data-race detector.

8 Configurations

Configurations implement a component specification by selecting regular components or instantiating generic components, and then connecting (“wiring”) these components together. The implementation section of a configuration has two parts:

```
configuration-implementation:
    implementation { component-listopt connection-list }
```

The *component-list* lists the components that are used to build this configuration, the *connection-list* specifies how these components are wired to each other and to the configuration’s specification. In the rest of this section, we call specification elements from the configuration’s specification *external*, and specification elements from one of the configuration’s components *internal*.

8.1 Included components

The *component-list* specifies the components used to build this configuration. These can be:

- A non-generic component *X*. Non-generic components are implicitly instantiated, references to *X* in different configurations will all refer to the same component.
- An instantiation of a generic component *Y*. Instantiations of *Y* in different configurations, or multiple instantiations in the same configuration will represent different components (see Section 6.1).

This *component-list* has the following syntax:

```
component-list:
    components
    component-list components

components:
    components component-line ;

component-line:
    component-ref instance-nameopt
    component-line , component-ref instance-nameopt

instance-name:
```

as identifier

component-ref:

identifier

new *identifier* (*component-argument-list*)

component-argument-list:

component-argument

component-argument-list , *component-argument*

component-argument:

expression

type-name

Each *component-ref* specifies a non-generic component *X* by simply giving its name (a compile-time error occurs if *X* is generic) and a generic component *Y* with **newY**(*args*) (a compile-time error occurs if *Y* is not generic). The arguments to *Y* must match the number of parameters *Y* expects, and:

- If the parameter is a type parameter, then the argument must be a non-function, non-array, non-incomplete type.
- If the parameter is of type **char**[], the argument must be a string constant.
- If the parameter is of arithmetic type, the argument must be a constant whose value is in the range of the parameter type.

Within the *connection-list*, a component specified in the *component-list* is referred to by:

- The name explicitly specified by the **X as Y** syntax (*instance-name*). This is necessary, e.g., when instantiating the same generic component more than once in a given configuration.
- The name of the component definition (**components new X()**, **Y**; is the same as **components new X() as X**, **Y as Y**).

The names specified by the *connection-list* belong to the object name space of the component's implementation scope (Section 4).

8.2 Wiring

Wiring is used to connect specification elements (interfaces, commands, events) together. This section and the next (Section 8.3) define the syntax and compile-time rules for wiring. Section 8.4 details how a program's wiring statements dictate which functions get called at each **call** and **signal** expression.

connection-list:

connection

connection-list connection

connection:

```
endpoint = endpoint  
endpoint -> endpoint  
endpoint <- endpoint
```

endpoint:

```
identifier-path  
identifier-path [ argument-expression-list ]
```

identifier-path:

```
identifier  
identifier-path . identifier
```

Wiring statements connect two *endpoints*. The *identifier-path* of an *endpoint* specifies a specification element. The *argument-expression-list* optionally specifies interface parameter values. We say that an endpoint is parameterised if its specification element is parameterised and the endpoint has no parameter values. A compile-time error occurs if an endpoint has parameter values and any of the following is true:

- The parameter values are not all constant expressions.
- The endpoint's specification element is not parameterised.
- There are more (or less) parameter values than there are parameters on the specification element.
- The parameter values are not in range for the specification element's parameter types.

A compile-time error occurs if the *identifier-path* of an *endpoint* is not of one the three following forms:

- X , where X names an external specification element.
- $K.X$ where K is a component from the *component-list* and X is a specification element of K .
- K where K is a some component name from the *component-list*. This form is used in implicit connections, discussed in Section 8.3. This form cannot be used when parameter values are specified.

Note that a component name can hide a specification element, preventing the element from being wired:

```
configuration AA { provides interface X as Y; }  
implementation {  
  components Z as Y, Z2 as Y2;
```

```

    Y /* refers to component Z, not interface X */ -> Y2.A;
}

```

Hiding specification elements will always result in a compile-time error as specification elements must all be wired.

There are three wiring statements in nesC:

- $endpoint_1 = endpoint_2$ (equate wires): Any connection involving an external specification element. These effectively make two specification elements equivalent.

Let S_1 be the specification element of $endpoint_1$ and S_2 that of $endpoint_2$. One of the following two conditions must hold or a compile-time error occurs:

- S_1 is internal, S_2 is external (or vice-versa) and S_1 and S_2 are both provided or both used,
- S_1 and S_2 are both external and one is provided and the other used.
- $endpoint_1 \rightarrow endpoint_2$ (link wires): A connection involving two internal specification elements. Link wires always connect a used specification element specified by $endpoint_1$ to a provided one specified by $endpoint_2$. If these two conditions do not hold, a compile-time error occurs.
- $endpoint_1 \leftarrow endpoint_2$ is equivalent to $endpoint_2 \rightarrow endpoint_1$.

In all three kinds of wiring, the two specification elements specified must be compatible, i.e., they must both be commands, or both be events, or both be interface instances. Also, if they are commands (or events), then they must both have the same function signature. If they are interface instances they must be of the same interface type. If these conditions do not hold, a compile-time error occurs.

If one endpoint is parameterised, the other must be too and must have the same parameter types; otherwise a compile-time error occurs.

The same specification element may be connected multiple times, e.g.:

```

configuration C {
    provides interface X;
} implementation {
    components C1, C2;

    X = C1.X;
    X = C2.X;
}

```

In this example, the multiple wiring will lead to multiple signalers (“fan-in”) for the events in interface X and for multiple functions being executed (“fan-out”) when commands in interface X are called. Note that multiple wiring can also happen when two configurations independently wire the same interface, e.g.:

<pre> configuration C { } implementation { components C1, C2; C1.Y -> C2.Y; } </pre>	<pre> configuration D { } implementation { components C3, C2; C3.Y -> C2.Y; } </pre>
----------------------------------------------------------------------------------------------	----------------------------------------------------------------------------------------------

All external specification elements must be wired or a compile-time error occurs. However, internal specification elements may be left unconnected (these may be wired in another configuration, or they may be left unwired if the modules have the appropriate `default` event or command implementations).

8.3 Implicit Connections

It is possible to write $K_1 \leftarrow K_2.X$ or $K_1.X \leftarrow K_2$ (and the same with $=$, or \rightarrow). This syntax iterates through the specification elements of K_1 (resp. K_2) to find a specification element Y such that $K_1.Y \leftarrow K_2.X$ (resp. $K_1.X \leftarrow K_2.Y$) forms a valid connection. If exactly one such Y can be found, then the connection is made, otherwise a compile-time error occurs.

For instance, with:

<pre> module M1 { provides interface StdControl; } ... </pre>	<pre> module M2 { uses interface StdControl as SC; } ... </pre>
-----------------------------------------------------------------	-------------------------------------------------------------------


```

configuration C { }
implementation {
  components M1, M2;
  M2.SC -> M1;
}

```

The `M2.SC -> M1` line is equivalent to `M2.SC -> M1.StdControl`.

8.4 Wiring Semantics

We first explain the semantics of wiring in the absence of parameterised interfaces. Section 8.4.1 below covers parameterised interfaces. Finally, Section 8.4.2 specifies requirements on the wiring statements of an application when viewed as a whole. We will use the simple application of Figure 1 as our running example.

For the purposes of this section, we will assume that all instantiations of generic components have been expanded into non-generic components as explained in Sections 6.1 and 10.2.

We define the meaning of wiring in terms of *intermediate functions*.⁴ There is one intermediate function I_α for every command or event α of every component. For instance, in Figure 1, module

⁴nesC can be compiled without explicit intermediate functions, so the behaviour described in this section has no runtime cost beyond the actual function calls and the runtime dispatch necessary for parameterised commands or events.

```

interface X {
  command int f();
  event void g(int x);
}

module M {
  provides interface X as P;
  uses interface X as U;
  provides command void h();
} implementation { ... }

configuration C {
  provides interface X;
  provides command void h2();
}

implementation {
  components M;
  X = M.P;
  M.U -> M.P;
  h2 = M.h;
}

```

Figure 1: Simple Wiring Example

M has intermediate functions $I_{M.P.f}$, $I_{M.P.g}$, $I_{M.U.f}$, $I_{M.U.g}$, $I_{M.h}$. In examples, we name intermediate functions based on their component, optional interface instance name and function name.

An intermediate function is either used or provided. Each intermediate function takes the same arguments as the corresponding command or event in the component's specification. The body of an intermediate function I is a list of calls (executed sequentially) to other intermediate functions. These other intermediate functions are the functions to which I is connected by the application's wiring statements. The arguments I receives are passed on to the called intermediate functions unchanged. The result of I is a list of results (the type of list elements is the result type of the command or event corresponding to I), built by concatenating the result lists of the called intermediate functions. An intermediate function which returns an empty result list corresponds to an unconnected command or event; an intermediate function which returns a list of two or more elements corresponds to "fan-out".

Intermediate Functions and Configurations The wiring statements in a configuration specify the body of intermediate functions. We first expand the wiring statements to refer to intermediate functions rather than specification elements, and we suppress the distinction between = and -> wiring statements. We write $I_1 \leftrightarrow I_2$ for a connection between intermediate functions I_1 and I_2 . For instance, configuration C from Figure 1 specifies the following intermediate function connections:

$$\begin{aligned}
I_{C.X.f} &\leftrightarrow I_{M.P.f} & I_{M.U.f} &\leftrightarrow I_{M.P.f} & I_{C.h2} &\leftrightarrow I_{M.h} \\
I_{C.X.g} &\leftrightarrow I_{M.P.g} & I_{M.U.g} &\leftrightarrow I_{M.P.g}
\end{aligned}$$

In a connection $I_1 \leftrightarrow I_2$ from a configuration C one of the two intermediate functions is the *callee* and the other is the *caller*. The connection simply specifies that a call to the callee is added to the body of the caller. I_1 (similarly, I_2) is a callee if any of the following conditions hold (we use the internal, external terminology for specification elements with respect to the configuration C containing the connection):

- If I_1 corresponds to an internal specification element that is a provided command or event.
- If I_1 corresponds to an external specification element that is a used command or event.
- If I_1 corresponds to a command of interface instance X , and X is an internal, provided or external, used specification element.
- If I_1 corresponds to an event of interface instance X , and X is an external, provided or internal, used specification element.

If none of these conditions hold, I_1 is a caller. The rules for wiring in Section 8.2 ensure that a connection $I_1 \leftrightarrow I_2$ cannot connect two callers or two callees. In configuration **C** from Figure 1, $I_{C.X.f}$, $I_{C.h2}$, $I_{M.P.g}$, $I_{M.U.f}$ are callers and $I_{C.X.g}$, $I_{M.P.f}$, $I_{M.U.g}$, $I_{M.h}$ are callees. Thus the connections of **C** specify that a call to $I_{M.P.f}$ is added to $I_{C.X.f}$, a call to $I_{C.X.g}$ is added to $I_{M.P.g}$, etc.

Intermediate Functions and Modules The **C** code in modules calls, and is called by, intermediate functions.

The intermediate function I for provided command or event α of module M contains a single call to the implementation of α in M . Its result is the singleton list of this call's result.

The expression `call $\alpha(e_1, \dots, e_n)$` is evaluated as follows:

- The arguments e_1, \dots, e_n are evaluated, giving values v_1, \dots, v_n .
- The intermediate function I corresponding to α is called with arguments v_1, \dots, v_n , with results list L .
- If $L = (w)$ (a singleton list), the result of the `call` is w .
- If $L = (w_1, w_2, \dots, w_m)$ (two or more elements), the result of the `call` depends on the result type τ of α . If $\tau = \text{void}$, then the result is `void`. Otherwise, τ must have an associated *combining function* c (Section 11.6 shows how combining functions are associated with types), or a compile-time error occurs. The combining function takes two values of type τ and returns a result of type τ . The result of the `call` is $c(w_1, c(w_2, \dots, c(w_{m-1}, w_m)))$ (note that the order of the elements of L was arbitrary).
- If L is empty the default implementation for α is called with arguments v_1, \dots, v_n , and its result is the result of the `call`. Section 8.4.2 specifies that a compile-time error occurs if L can be empty and there is no default implementation for α .

The rules for `signal` expressions are identical.

Example Intermediate Functions Figure 2 shows the intermediate functions that are produced for the components of Figure 1, using a C-like syntax, where `list(x)` produces a singleton list containing x , `empty_list` is a constant for the 0 element list and `concat_list` concatenates two lists. The calls to `M.P.f`, `M.U.g`, `M.h` represent calls to the command and event implementations in module **M** (not shown).

list of int $I_{M.P.f}()$ { return list(M.P.f()); }	list of void $I_{M.P.g}(\text{int } x)$ { list of int $r1 = I_{C.X.g}(x)$; list of int $r1 = I_{M.U.g}(x)$; return list_concat(r1, r2); }
list of int $I_{M.U.f}()$ { return $I_{M.P.f}()$; }	list of void $I_{M.U.g}(\text{int } x)$ { return list(M.U.g(x)); }
list of int $I_{C.X.f}()$ { return $I_{M.P.f}()$; }	list of void $I_{C.X.g}(\text{int } x)$ { return empty_list; }
list of void $I_{C.h2}()$ { return $I_{M.h}()$; }	list of void $I_{M.h}()$ { return list(M.h()); }

Figure 2: Intermediate Functions for Figure 1

8.4.1 Wiring and Parameterised Functions

If a command or event α of component K is parameterised with interface parameters of type τ_1, \dots, τ_n then there is an intermediate function $I_{\alpha, v_1, \dots, v_n}$ for every distinct tuple $(v_1 : \tau_1, \dots, v_n : \tau_n)$.

In modules, if intermediate function I_{v_1, \dots, v_n} corresponds to parameterised, provided command (or event) α then the call in I_{v_1, \dots, v_n} to α 's implementation passes values v_1, \dots, v_n as the values for α 's interface parameters.

The expression `call $\alpha[e'_1, \dots, e'_m](e_1, \dots, e_n)$` is evaluated as follows:

- The arguments e_1, \dots, e_n are evaluated, giving values v_1, \dots, v_n .
- The arguments e'_1, \dots, e'_m are evaluated, giving values v'_1, \dots, v'_m .
- The v'_i values are cast to type τ_i , where τ_i is the type of the i th interface parameter of α .
- The intermediate function $I_{v'_1, \dots, v'_m}$ corresponding to α is called with arguments v_1, \dots, v_n , with results list L .⁵
- If L has one or more elements, the result of the `call` is produced as in the non-parameterised case.
- If L is empty the default implementation for α is called with interface parameter values v'_1, \dots, v'_m and arguments v_1, \dots, v_n , and its result is the result of the `call`. Section 8.4.2

⁵This call typically involves a runtime selection between several command implementations - this is the only place where intermediate functions have a runtime cost.

specifies that a compile-time error occurs if L can be empty and there is no default implementation for α .

The rules for **signal** expressions are identical.

There are two cases when an endpoint in a wiring statement refers to a parameterised specification element:

- The endpoint specifies parameter values v_1, \dots, v_n . If the endpoint corresponds to commands or events $\alpha_1, \dots, \alpha_m$ then the corresponding intermediate functions are $I_{\alpha_1, v_1, \dots, v_n}, \dots, I_{\alpha_m, v_1, \dots, v_n}$ and wiring behaves as before.
- The endpoint does not specify parameter values. In this case, both endpoints in the wiring statement correspond to parameterised specification elements, with identical interface parameter types τ_1, \dots, τ_n . If one endpoint corresponds to commands or events $\alpha_1, \dots, \alpha_m$ and the other to corresponds to commands or events β_1, \dots, β_m , then there is a connection $I_{\alpha_i, w_1, \dots, w_n} \leftrightarrow I_{\beta_i, w_1, \dots, w_n}$ for all $1 \leq i \leq m$ and all tuples $(w_1 : \tau_1, \dots, w_n : \tau_n)$ (i.e., the endpoints are connected for all corresponding parameter values).

8.4.2 Application-level Requirements

There are two requirement that the wiring statements of an application must satisfy, or a compile-time error occurs:

- There must be no infinite loop involving only intermediate functions.
- At every **call** α (or **signal** α) expression in the application's modules:
 - If the call is unparameterised: if the call returns an empty result list there must be a default implementation of α (the number of elements in the result list depends only on the wiring).
 - If the call is parameterised: if substitution of any values for the interface parameters of α returns an empty result list there must be a default implementation of α (the number of elements in the result list for a given parameter value tuple depends only on the wiring).

Note that this condition does not consider the expressions used to specify interface parameter values at the call-site.

9 Concurrency in nesC

nesC assumes an execution model that consists of run-to-completion *tasks* (that typically represent the ongoing computation), and *interrupt handlers* that are signaled asynchronously by hardware. A scheduler for nesC can execute tasks in any order, but must obey the run-to-completion rule (the standard TinyOS scheduler follows a FIFO policy). Because tasks are not preempted and run to completion, they are atomic with respect to each other, but are not atomic with respect to interrupt handlers.

As this is a concurrent execution model, nesC programs are susceptible to race conditions, in particular data races on the program’s *shared state*, i.e., its global and module variables (nesC does not include dynamic memory allocation). Races are avoided either by accessing a shared state only in tasks, or only within atomic statements. The nesC compiler reports potential data races to the programmer at compile-time.

Formally, we divide the code of a nesC program into two parts:

Synchronous Code (SC): code (functions, commands, events, tasks) that is only reachable from tasks.

Asynchronous Code (AC): code that is reachable from at least one interrupt handler.

Although non-preemption eliminates data races among tasks, there are still potential races between SC and AC, as well as between AC and AC. In general, any update to shared state that is *reachable from AC* is a potential data race. The basic invariant nesC enforces is:

Race-Free Invariant: *Any update to shared state is either SC-only or occurs in an atomic statement.* The body of a function f called from an atomic statement is considered to be “in” the atomic statement as long as all calls to f are “in” atomic statements.

It is possible to introduce a race condition that the compiler cannot detect, but it must span multiple atomic statements or tasks and use storage in intermediate variables.

nesC may report data races that cannot occur in practice, e.g., if all accesses are protected by guards on some other variable. To avoid redundant messages in this case, the programmer can annotate a variable v with the `norace` storage-class specifier to eliminate all data race warnings for v . The `norace` keyword should be used with caution.

nesC reports a compile-time error for any command or event that is AC and that was not declared with `async`. This ensures that code that was not written to execute safely in an interrupt handler is not called inadvertently.

10 nesC Applications

A nesC application has three parts: a list of C declarations and definitions, a set of interface definitions and a set of components (non-generic components and instantiated generic components).

The C declarations and definitions, interfaces and components that form a nesC application are determined by an on-demand loading process. The input to the nesC compiler is a single non-generic component K . The nesC compiler first loads a user-specified set of C files⁶ (Section 10.1), then loads component K (Section 10.2). The code for the application is all the C code loaded and all the components generated as a result of loading these two files. A nesC compiler can assume that all calls to functions, commands or events not marked with the `spontaneous`, `hwevent` or `atomic.hwevent` attributes (Section 11.6) occur in the loaded code (i.e., there are no “invisible” calls to any other functions).⁷

⁶`ncc`, the TinyOS frontend for nesC always loads the TinyOS `tos.h` file.

⁷For instance, the current nesC compiler uses this information to eliminate unreachable code.

During preprocessing of loaded files, nesC defines the `NESC` symbol to a number XYZ which identifies the version of the nesC language and compiler. For nesC 1.1, XYZ is at least 110.⁸

Part of the process of loading a C file, nesC component or interface definition involves locating the corresponding source file. The mechanism used to locate files is outside the scope of this reference manual; for details on how this works in the current compiler please see the `ncc` man page.

10.1 Loading C file X

If X has already been loaded, nothing more is done. Otherwise, file $X.h$ is located and preprocessed. Changes made to C macros (via `#define` and `#undef`) are visible to all subsequently preprocessed files. The C declarations and definitions from the preprocessed $X.h$ file are entered into the C global scope, and are therefore visible to all subsequently processed C files, interfaces and components.

10.2 Loading Component K

If K has already been loaded, nothing more is done. Otherwise, file $K.nc$ is located and preprocessed. Changes made to C macros (via `#define` and `#undef`) are discarded. The preprocessed file is parsed using the following grammar:

```

nesC-file:
    includes-listopt interface
    includes-listopt module
    includes-listopt configuration

includes-list:
    includes
    includes-list includes

includes:
    includes identifier-list ;

```

If $K.nc$ does not define `module` K or `configuration` K , a compile-time error is reported. Otherwise, all C files specified by the *includes-list* are loaded (Section 10.1 possibly). Note that as file $K.nc$ has already been loaded, macros defined in these files cannot be used in $K.nc$ (though they will be visible for subsequently loaded files).

Next, all interface definitions used in the component's specification are loaded (Section 10.3). Next, the component specification is processed (Section 5.3). If K is a configuration, all component definitions used (Section 8.1) in K are loaded (Section 10.2). The implementation of K is then checked according to the rules of Section 7 or Section 8.

Finally, K is processed to produce a set of output components via the following algorithm:

Expand(K):

⁸The `NESC` symbol was not defined in earlier versions of nesC.

- If K is a generic component, no component is output.
- If K is a non-generic module, K is output.
- If K is a non-generic configuration: for each component instantiation `newL(a1, ..., an)` in K , a new component X is created according to the rules of Section 6.1 and `Expand(X)` is called recursively (instantiating further generic components if L contained component instantiations). Then K is output.

10.3 Loading Interface Definition I

If I has already been loaded, nothing more is done. Otherwise, file $I.nc$ is located and preprocessed. Changes made to C macros (via `#define` and `#undef`) are discarded. The preprocessed file is parsed following the *nesC-file* production above. If $I.nc$ does not define `interface I` a compile-time error is reported. Otherwise, all C files specified by the *includes-list* are loaded (Section 10.1) (as with components, macros defined in these C files cannot be used in $I.nc$ as it has already been preprocessed). Then I 's definition is processed (Section 5).

As an example of including C files in components or interfaces, interface definition `Bar` might include C file `BarTypes.h` which defines types used in `Bar`:

<pre>Bar.nc: includes BarTypes; interface Bar { command result_t bar(BarType arg1); }</pre>	<pre>BarTypes.h: typedef struct { int x; double y; } BarType;</pre>
-------------------------------------------------------------------------------------------------	-----------------------------------------------------------------------------

The definition of interface `Bar` can refer to `BarType`, as can any component that uses or provides interface `Bar` (interface `Bar`, and hence `BarTypes.h`, are loaded before any such component's specification or implementation are processed).

11 Miscellaneous

11.1 Constant Folding in nesC

There are two extensions to C's constant folding (see A.7.19 in [2]) in nesC: *constant functions* and *unknown constants*. *Constant Functions* are functions provided by the nesC language which return a compile-time constant. The definition of nesC's constant functions is given in Section 11.2. An *unknown constant* is a constant whose value is not yet known, e.g., non-type parameters to generic components are unknown constants. Unknown constants allow the definition of a generic component to be (mostly, see below) checked for correctness before its arguments' values are known.

An expression involving an unknown constant is considered a constant expression if the resulting expression is constant irrespective of the unknown constant's value, with the following exceptions: a/b and $a\%b$ can assume that b is not zero. Constant expressions involving unknown constants

are re-checked once the values of constant expressions become known.⁹ As a result, the following generic component definition is legal:

```
generic module A(int n) { }
implementation {
    int s = 20 / n;
}
```

but the following instantiation will report a compile-time error:

```
configuration B { }
implementation {
    components new A(0) as MyA;
}
```

11.2 Compile-time Constant Functions

nesC currently has two constant functions:

- `unsigned int unique(char *identifier)`

Returns: each call to `unique` with the same `identifier` string returns a different value. The expansion of `unique` calls happens after generic component instantiation (Section 10.2): calls to `unique` in generic components return a different value in each instantiation. The range of values returned by `unique(s)` is $0..n - 1$, where n is the number of different values returned by `unique(s)`. For purposes of checking constant expressions, `unique(s)` behaves as if it were an unknown constant.

The intended use of `unique` is for passing a unique integer to parameterised interface instances, so that a component providing a parameterised interface can uniquely identify the various components connected to that interface.

- `unsigned int uniqueCount(char *identifier)`

Returns: `uniqueCount(s)` returns one more than the largest value returned by any of the calls to `unique(s)`, or 0 if there are no calls to `unique(s)`. For purposes of checking constant expressions, `uniqueCount(s)` behaves as if it were an unknown constant.

The intended use of `uniqueCount` is for dimensioning arrays (or other data structures) which will be indexed using the numbers returned by `unique`. For instance, a `Timer` service that identifies its clients (and hence each independent timer) via a parameterised interface and `unique` can use `uniqueCount` to allocate the correct number of timer data structures.

In the following example:

```
generic module A() { }
implementation {
    int x = unique("A");
}
```

⁹The time at which the value of unknown constants become known is unspecified by this language definition.

```

    int y = uniqueCount("A");
}
configuration B {
implementation {
    components new A() as A1, new A() as A2;
}

```

$B.A1.y = B.A2.y = 2$ and either $B.A1.x = 0$, $B.A2.x = 1$ or $B.A1.x = 1$, $B.A2.x = 0$.

11.3 Type Parameters and C Type Checking

Generic interface and component definitions can have type parameters. Syntactically, type parameters behave the same as `typedef`'d identifiers. When a generic component or interface is instantiated, the type parameter will be replaced with the argument type, which cannot be incomplete, of function or of array type. The size and alignment of a type parameter are an unknown constant (Section 11.1). The rules for assignment and type equivalence for a type parameter t are simple: a value of type t is assignable to an lvalue of type t (extends A.7.17 in [2]) and type t is only equivalent to itself (extends A.8.10 in [2]).

- To build new

11.4 Functions with no arguments, old-style C declarations

nesC functions with no arguments are declared with `()`, not `(void)`. The latter syntax reports a compile-time error.

Old-style C declarations (with `()`) and function definitions (parameters specified after the argument list) are not allowed in interfaces or components (and cause compile-time errors).

Note that neither of these changes apply to C files (so that existing `.h` files can be used unchanged).

11.5 `//` comments

nesC allows `//` comments in C, interface and component files.

11.6 Attributes

nesC uses gcc's¹⁰ `__attribute__` syntax for declaring some properties of functions, variables and typedefs. These attributes can be placed either on declarations (after the declarator) or function definitions (after the parameter list).¹¹ The attributes of x are the union of all attributes on all declarations and definitions of x .

The attribute syntax in nesC is:

¹⁰<http://gcc.gnu.org>

¹¹gcc doesn't allow attributes after the parameter list in function definitions.

init-declarator-list: also

init-declarator attributes

init-declarator-list , *init-declarator attributes*

function-definition: also

*declaration-specifiers*_{opt} *declarator attributes* *declaration-list*_{opt} *compound-statement*

attributes:

attribute

attributes attribute

attribute:

`__attribute__` ((*attribute-list*))

attribute-list:

single-attribute

attribute-list , *single-attribute*

single-attribute:

identifier

identifier (*argument-expression-list*)

nesC supports three attributes:

- **C**: This attribute is used for a C declaration or definition *d* at the top-level of a module (it is ignored for all other declarations). It specifies that *d*'s should appear in the global C scope rather than in the module's per-component-implementation scope. This allows *d* to be used (e.g., called if it is a function) from C code.
- **spontaneous**: This attribute can be used on any function *f* (in modules or C code). It indicates that there are calls *f* that are not visible in the source code. Typically, functions that are called spontaneously are interrupt handlers, and the C `main` function. Section 10 discusses how the nesC compiler uses the **spontaneous** attribute during compilation.
- **hwevent**: This attribute can be used on any function *f* (in modules or C code). It indicates that *f* is an interrupt handler, i.e., that there are spontaneous calls to *f* and that *f* is AC (Section 9).
- **atomic_hwevent**: This attribute can be used on any function *f* (in modules or C code). This behaves the same as **hwevent**, but, additionally, informs nesC that the body of *f* behaves as if it were an **atomic** statement (on typical hardware this means that this interrupt handler runs with interrupts disabled).

Note that neither **hwevent** or **atomic_hwevent** provide any linkage of *f* with a particular interrupt handler. The mechanism by which that is achieved is platform-specific.

- `combine(fname)`: This attribute specifies the combining function for a type in a `typedef` declaration. The combining function specifies how to combine the multiple results of a call to a command or event which has “fan-out”. For example:

```
typedef uint8_t result_t __attribute__((combine(rcombine)));

result_t rcombine(result_t r1, result_t r2)
{
    return r1 == FAIL ? FAIL : r2;
}
```

specifies logical-and-like behaviour when combining commands (or events) whose result type is `result_t`. See Section 8.4 for the detailed semantics.

A compile-time error occurs if the combining function c for a type t does not have the following type: $t\ c(t, t)$.

Example of attribute use: in file `RealMain.td`:

```
module RealMain { ... }
implementation {
    int main(int argc, char **argv) __attribute__((C, spontaneous)) {
        ...
    }
}
```

This example declares that function `main` should actually appear in the C global scope (C), so that the linker can find it. It also declares that `main` can be called even though there are no function calls to `main` anywhere in the program (`spontaneous`).

A Grammar

Please refer to Appendix A of Kernighan and Ritchie (K&R) [2, pp234–239] while reading this grammar.

The following keywords are new for nesC: `as`, `call`, `command`, `components`, `configuration`, `event`, `implementation`, `interface`, `module`, `post`, `provides`, `signal`, `task`, `uses`, `includes`. These nesC keywords are not reserved in C files. The corresponding C symbols are accessible in nesC files by prefixing them with `__nesc_keyword` (e.g., `__nesc_keyword_as`).

nesC reserves all identifiers starting with `__nesc` for internal use. TinyOS reserves all identifiers starting with `TOS_` and `TOSH_`.

nesC files follow the *nesC-file* production; `.h` files included via the `includes` directive follow the *translation-unit* directive from K&R.

New rules:

nesC-file:

- includes-list_{opt} interface*
- includes-list_{opt} module*
- includes-list_{opt} configuration*

includes-list:

- includes*
- includes-list includes*

includes:

- includes** *identifier-list* ;

interface:

- interface** *identifier* { *declaration-list* }

module:

- module** *identifier specification module-implementation*

module-implementation:

- implementation** { *translation-unit* }

configuration:

- configuration** *identifier specification configuration-implementation*

configuration-implementation:

- implementation** { *component-list_{opt} connection-list* }

component-list:

- components*
- component-list components*

components:

- components** *component-line* ;

component-line:

- renamed-identifier*
- component-line* , *renamed-identifier*

renamed-identifier:

- identifier*
- identifier as identifier*

connection-list:

- connection*

connection-list connection

connection:

endpoint = endpoint
endpoint -> endpoint
endpoint <- endpoint

endpoint:

identifier-path
identifier-path [argument-expression-list]

identifier-path:

identifier
identifier-path . identifier

specification:

{ uses-provides-list }

uses-provides-list:

uses-provides
uses-provides-list uses-provides

uses-provides:

uses *specification-element-list*
provides *specification-element-list*

specification-element-list:

specification-element
{ specification-elements }

specification-elements:

specification-element
specification-elements specification-element

specification-element:

declaration
interface *renamed-identifier parameters_{opt}*

parameters:

[parameter-type-list]

Changed rules:

storage-class-specifier: also one of

command event async task norace

declaration-specifiers: also
 default *declaration-specifiers*

direct-declarator: also
 identifier . *identifier*
 direct-declarator *parameters* (*parameter-type-list*)

init-declarator-list: also
 init-declarator *attributes*
 init-declarator-list , *init-declarator* *attributes*

function-definition: also
 *declaration-specifiers*_{opt} *declarator* *attributes* *declaration-list*_{opt} *compound-statement*

attributes:
 attribute
 attributes *attribute*

attribute:
 __attribute__ ((*attribute-list*))

attribute-list:
 single-attribute
 attribute-list , *single-attribute*

single-attribute:
 identifier
 identifier (*argument-expression-list*)

statement: also
 atomic-statement

atomic-statement:
 atomic *statement*

postfix-expression: replaced by
 primary-expression
 postfix-expression [*argument-expression-list*]
 *call-kind*_{opt} *primary* (*argument-expression-list*_{opt})
 postfix-expression . *identifier*
 postfix-expression -> *identifier*
 postfix-expression ++
 postfix-expression --

call-kind: one of
call signal post

B Glossary

- *combining function*: C function that combines the multiple results of command call (or event signal) in the presence of *fan-out*.

- *command, event*: A function that is part of a component's *specification*, either directly as a *specification element* or within one of the component's *interface instances*.

When used directly as specification elements, commands and events have roles (*provider, user*) and can have *interface parameters*. As with interface instances, we distinguish between *simple commands (events)* without interface parameters and *parameterised commands (events)* with interface parameters. The interface parameters of a command or event are distinct from its regular function parameters.

- *compile-time error*: An error that the nesC compiler must report at compile-time.
- *component*: The basic unit of nesC programs. Components have a name and are of two kinds: *modules* and *configurations*. A component has a *specification* and an implementation.
- *configuration*: A component whose implementation is provided by a composition of other components with a specific *wiring*.
- *endpoint*: A specification of a particular specification element, and optionally some interface parameter values, in a wiring statement of a configuration. A parameterised endpoint is an endpoint without parameter values that corresponds to a parameterised specification element.
- *event*: See command.
- *extent*: The lifetime of a variable. nesC has the standard C extents: *indefinite, function, and block*.
- *external*: In a configuration *C*, describes a specification element from *C*'s specification. See internal.
- *fan-in*: Describes a provided command or event called from more than one place.
- *fan-out*: Describes a used command or event connected to more than one command or event implementation. A *combining function* combines the results of calls to these used commands or events.
- *interface*: When the context is unambiguous, we use interface to refer to either an *interface type* or an *interface instance*.
- *interface instance*: An instance of a particular *interface type* in the *specification* of a component. An interface instance has an instance name, a role (*provider* or *user*), an *interface type* and, optionally, *interface parameters*. An interface instance without parameters is a *simple interface instance*, with parameters it is a *parameterised interface instance*.

- *interface parameter*: An interface parameter has an interface parameter name and must be of integral type.

There is (conceptually) a separate *simple interface instance* for each distinct list of parameter values of a *parameterised interface instance* (and, similarly, separate simple commands or events in the case of parameterised commands or events). Parameterised interface instances allow runtime selection based on parameter values between a set of commands (or between a set of events).

- *interface type*: An *interface type* specifies the interaction between two components, *the provider* and the *user*. This specification takes the form of a set of *commands* and *events*. Each interface type has a distinct name.

Interfaces are bi-directional: the provider of an interface implements its commands, the user of an interface implements its events.

- *intermediate function*: A pseudo-function that represents the behaviour of the commands and events of a component, as specified by the wiring statements of the whole application. See Section 8.4.
- *internal*: In a configuration *C*, describes a specification element from one of the components specified in *C*'s component list. See external.
- *module*: A component whose implementation is provided by C code.

- *namespace*: nesC has the standard C *variable* (also used for functions and **typedefs**), *tagged type* (**struct**, **union** and **enum** tag names) and *label* namespaces. Additionally, nesC has a *component and interface type* namespace for component and interface type names.

parameterised command, *parameterised event*, *parameterised interface instance*, *endpoint*: See command, event, interface instance, endpoint.

provided, *provider*: A role for a specification element. Providers of *interface instances* must implement the *commands* in the interface; provided commands and events must be implemented.

provided command of K: A command that is either a provided specification element of *K*, or a command of a provided interface of *K*.

provided event of K: An event that is either a provided specification element of *K*, or an event of a used interface of *K*.

- *scope*: nesC has the standard C *global*, *function-parameter* and *block* scopes. Additionally there are *specification* and *implementation* scopes in components and a per-interface-type scope. Scopes are divided into namespaces.

simple command, *simple event*, *simple interface instance*: See command, event, interface instance.

- *specification*: A list of *specification elements* that specifies the interaction of a component with other components.
- *specification element*: An *interface instance*, *command* or *event* in a specification that is either *provided* or *used*.

- *task*: A TinyOS task.
used, user: A role for a specification element. Users of *interface instances* must implement the *events* in the interface.
used command of K: A command that is either a used specification element of *K*, or a command of a used interface of *K*.
used event of K: An event that is either a used specification element of *K*, or an event of a provided interface of *K*.
- *wiring*: The connections between component's specification elements specified by a configuration.

References

- [1] J. Hill, R. Szewczyk, A. Woo, S. Hollar, D. E. Culler, and K. S. J. Pister. System Architecture Directions for Networked Sensors. In *Architectural Support for Programming Languages and Operating Systems*, pages 93–104, 2000. TinyOS is available at <http://webs.cs.berkeley.edu>.
- [2] B. W. Kernighan and D. M. Ritchie. *The C Programming Language, Second Edition*. Prentice Hall, 1988.